

–Appendix–

Prune but Preserve: Decoupling Redundancy Suppression and Semantic Preservation for Efficient Time Series Clustering

Zexi Tan¹, Haoyi Xiao¹, Linhao Ouyang¹, Mengke Li², Yiqun Zhang¹ and Yang Lu³

¹Guangdong University of Technology, ²Shenzhen University, ³Xiamen University
yqzhang@gdut.edu.cn, mengkeli@szu.edu.cn, luyang@xmu.edu.cn

I. BASELINE DESCRIPTIONS

The baseline set is organized around the empirical question studied in the main paper, namely at which stage redundancy should be reduced for unsupervised MTS clustering. It covers five streams: full-sequence representation learning, input-side token exclusion, partial-observation clustering, structural subsequence indexing, and foundation-style time series modeling. This design separates the benefit of source-level patch admission from gains brought by contrastive objectives, spectral modeling, graph structure, incomplete-view learning, or general temporal pretraining. For embedding-output baselines, k -means gives final assignments with the label-provided cluster number k , following the same unsupervised protocol as the main paper.

FCACC [1] is a fuzzy cluster-aware contrastive clustering method for time series. It constructs augmented views and introduces cluster-aware hard negatives to reduce the gap between instance contrast and cluster assignment. FCACC represents a clustering-oriented full-sequence baseline, testing whether cluster-aware contrast can reduce background influence without source-level patch admission. The official implementation is available at [FCACC](#).

TFMCC [2] performs time-frequency augmented multi-level contrastive clustering. Temporal and spectral augmentations enhance cluster-relevant patterns at different representation levels, but the complete sequence still enters the encoding pipeline. TFMCC tests whether temporal-spectral contrast can handle redundancy as effectively as pruning before encoder-side interaction. The official implementation is available at [TFMCC](#).

GTM [3] learns general time series representations through frequency-domain attention and shared temporal modeling. As GTM is designed for transferable representation rather than clustering-specific reduction, k -means is applied to its embeddings for final assignment. The comparison examines whether broad temporal pretraining can replace explicit patch admission under redundancy-dominated MTS records. The official implementation is available at [GTM](#).

EMTC [4] is the closest input-side exclusion baseline for

MTS clustering. It suppresses redundancy through importance-aware variate-wise masking and enriches retained content with multi-endogenous views. Its comparison with SIGMA directly tests whether confidence-calibrated admission, routed semantic restoration, and prototype-guided refinement improve clustering after sparse token admission. The official implementation is available at [EMTC](#).

FEI [5] introduces a frequency-domain information bottleneck by masking frequency components and predicting the corresponding embeddings. This design encourages informative spectral structure without manually defined positive and negative pairs. FEI serves as a spectral-reduction baseline, assessing whether patch-level admission still helps beyond frequency-oriented representation learning. The official implementation is available at [FEI](#).

MVCIMTS [6] is a contrastive multi-view clustering method for incomplete MTS. It treats variates as views and learns complementary information through multi-view contrast. Since its partial observation comes from variables instead of temporal patches, MVCIMTS tests whether variable-wise incompleteness brings the same clustering benefit as temporal admission followed by routed restoration. The official implementation is available at [MVCIMTS](#).

k -Graph [7] converts subsequences into graph representations and clusters graph-derived features. By building graphs across multiple subsequence lengths, k -Graph captures recurring local patterns and structural similarity. It provides a data-centric baseline for datasets whose clusters depend on shape correspondence or repeated subsequence structure. The official implementation is available at [k-Graph](#).

TimesURL [8] is a universal self-supervised representation method for time series. It combines contrastive learning with temporal constraints to obtain general embeddings. In the prune-and-preserve comparison, TimesURL represents a full-sequence self-supervised baseline because redundancy is not controlled through an admissible patch set. The official implementation is available at [TimesURL](#).

UNITS [9] is a unified multi-task time series model. It represents task information as tokens in a shared modeling framework and learns transferable representations across domains,

temporal scales, and sampling patterns. UNITS provides a foundation-style baseline for testing whether task-unified pre-training is sufficient for redundancy-dominated MTS clustering. Its comparison with SIGMA highlights the role of source-level admission and routed semantic preservation in cluster-oriented representation learning. The official implementation is available at [UNITS](#).

II. IMPLEMENTATION DETAILS

All experiments follow the unsupervised clustering protocol in the main paper. For each UEA dataset, training and test splits are concatenated before clustering, and labels are accessed only for metric computation. Each experiment is repeated with five random seeds, and the mean and standard deviation are reported. This protocol keeps the appendix consistent with the main-paper evaluation and makes the supplementary tables comparable to the representative results in Section IV of the full paper.

SIGMA uses Confidence-Calibrated Masked Admission as the default admission strategy. In the masking ablation, source-level admission is removed while all other settings remain unchanged, isolating the effect of controlling patch participation before encoder-side feature interaction. Unless otherwise specified, no input perturbation is added, and Gaussian noise is introduced only in the robustness study.

For patch-level admission, the patch length is set to $\ell_p = 4$. Given $m = \lceil T/\ell_p \rceil$ patches, the retained ratio is $r_a = 0.5$, and the retained patch number is $m_a = \lceil r_a m \rceil$. Thus, the encoder processes a compact patch subset after admission. This setting reduces redundant patch participation before feature interaction while keeping enough temporal context for semantic restoration.

For endogenous-view representation learning, the number of views is $n_v = 2$. Each view encoder uses hidden dimension $d_h = 32$, and the fused clustering embedding has dimension $d_z = 64$. The prototype alignment temperature is $\tau_p = 0.1$. These values give SIGMA a compact representation budget while allowing admitted content to be organized into complementary latent views.

SIGMA is trained for $e_{\max} = 200$ epochs with mini-batch size $b = 64$. Adam is adopted with learning rate $\alpha = 0.001$. The loss weights are $\lambda_{\text{rec}} = 1.0$, $\lambda_{\text{orth}} = 0.1$, and $\lambda_{\text{pa}} = 1.0$, corresponding to reconstruction, inter-view orthogonality, and prototype alignment, respectively. Reconstruction anchors temporal content after admission, orthogonality discourages view homogenization, and prototype alignment sharpens the fused cluster space. Under this configuration, admission control, semantic restoration, and cluster refinement keep distinct roles within the prune-and-preserve pipeline.

III. SUPPLEMENTARY RESULTS

This section provides the supplementary empirical record behind the main paper. The main paper keeps the evaluation compact through representative datasets and condensed diagnostics. The appendix reports complete benchmark tables, Bonferroni-Dunn rank tests, full module and route ablations,

additional training-loss curves, mask-rate scalability, and the robustness and compactness diagnostics in Fig. 4 under the same unsupervised clustering protocol defined in Section II. These results do not introduce another evaluation setting. They make the main claims traceable at the dataset, component, optimization, statistical, and resource levels.

The organization follows the role assignment in SIGMA. Complete clustering results and rank tests examine whether source-level admission and routed representation learning remain competitive across heterogeneous MTS collections. Ablations verify the contributions of admission, endogenous-view organization, semantic restoration, and prototype-guided refinement. Loss curves inspect whether staged training stays controlled after prototype guidance is introduced. Mask-rate scalability checks whether runtime reduction comes from fewer admitted patches before encoder-side interaction. Robustness and compactness diagnostics further test whether the retained representation remains usable under noise perturbation and reduced embedding dimensions. Together, these supplementary results connect clustering quality, statistical reliability, mechanism attribution, training stability, and computation-aware temporal data organization.

A. Complete Clustering Results on 30 UEA Datasets

The main paper answers Q1 with representative per-dataset results and a Bonferroni-Dunn rank test. This appendix subsection provides the complete dataset-level record. TABLE I and TABLE II report ACC, F1, NMI, and ARI on all 30 UEA datasets, covering both the datasets discussed in the main paper and the remaining benchmarks omitted due to page limits. The appendix therefore supports Q1 with the full benchmark evidence under the same unsupervised protocol.

The complete tables show that SIGMA’s advantage does not come from a few favorable cases. On datasets with clear cluster-revealing patterns, such as ArticulatoryWordRecognition, BasicMotions, JapaneseVowels, PenDigits, and SelfRegulationSCP1, SIGMA obtains leading or highly competitive scores across multiple metrics. These datasets are important for the prune-and-preserve claim, because sparse class-related changes can be diluted when complete sequences enter feature interaction. By reducing background-dominated patch participation before encoding and organizing retained content through routed objectives, SIGMA keeps compact representations informative for clustering. JapaneseVowels and PenDigits are typical cases where reduced input participation still yields a discriminative grouping space.

The full benchmark also reflects the diversity of unsupervised MTS clustering. Several datasets favor specialized baselines when cluster similarity depends more on shape matching, spectral structure, or weak label-to-signal correspondence. For example, FCACC, TFMCC, EMTC, FEI, and k -Graph obtain the best result on some individual datasets. Thus, the key observation is not dominance on every dataset, but stable average behavior across different sample sizes, sequence lengths, variable dimensions, and label-provided cluster numbers k . Fig. 1 provides the corresponding Bonferroni-Dunn




Fig. 2. Additional training-loss curves during SIGMA training on six UEA datasets. Each curve reports the mean loss across five random seeds, and the shaded band denotes standard deviation. Vertical dashed lines mark the semantic warm-up, dual-route ramp-up, and prototype-guided refinement stages.




Fig. 3. Scalability under increasing mask rates. The three panels vary N , T , and D , respectively. For each setting, execution time is normalized by the cost at zero mask rate. Shaded bands denote standard deviation, and endpoint annotations report the reduction at 90% mask rate.




Fig. 4. **Robustness and compactness diagnostics.** F1-score on AtrialFibrillation and Libras under (a) increasing Gaussian noise and (b) reduced embedding dimensions from 64 to 4. SIGMA keeps stable accuracy in both settings, showing that retained representations remain usable under perturbation and compact latent budgets.

as the noise level increases, especially on AtrialFibrillation, where the curve stays above the compared baselines. On Libras, several methods fluctuate or decline after perturbation, whereas SIGMA remains stable across noise levels. The result suggests that confidence-calibrated admission filters part of the corrupted or routine content before feature interaction, and the retained views still provide sufficient context for clustering.

Fig. 4(b) further reduces the embedding dimension from 64 to 4. SIGMA maintains competitive F1-score under compressed dimensions on both datasets, showing that the learned cluster structure does not rely on an oversized latent space. EMTC remains competitive on Libras at the smallest dimension, but SIGMA gives steadier behavior across the full dimension range. Together, the two diagnostics indicate that the prune-but-preserve design supports perturbation tolerance and compact representation learning, which is essential for efficient MTS organization under constrained computation and storage budgets.

REFERENCES

- [1] C. Wang, M. Du, X. Jiang, and Y. Dong, “Fuzzy cluster-aware contrastive clustering for time series,” *Pattern Recognition*, vol. 173, p. 112899, 2026.
- [2] C. Wang, M. Du, and X. Jiang, “Time-frequency augmented multi-level contrastive clustering for time series,” in *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, vol. 40, no. 31, 2026, pp. 26 142–26 150.
- [3] C. HE, X. Huang, G. Jiang, Z. Li, D. Lian, H. Xie, E. Chen, xijie liang, Zhengzengrong, and P. Lee, “GTM: A general time-series model for enhanced representation learning of time-series data,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2026. [Online]. Available: <https://openreview.net/forum?id=PWM6FERWz9>
- [4] Z. Tan, X. Luo, Y. Liu, and Y. Zhang, “Mask the redundancy: Evolving masking representation learning for multivariate time-series clustering,” in *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, vol. 40, no. 30, 2026, pp. 25 787–25 795.
- [5] E. Fu and Y. Hu, “Frequency-masked embedding inference: A non-contrastive approach for time series representation learning,” in *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, vol. 39, no. 16, 2025, pp. 16 639–16 647.
- [6] Y. Li, M. Du, X. Jiang, and N. Zhang, “Contrastive learning-based multi-view clustering for incomplete multivariate time series,” *Information Fusion*, vol. 117, p. 102812, 2025.
- [7] P. Boniol, D. Tiano, A. Bonifati, and T. Palpanas, “ k -graph: A graph embedding for interpretable time series clustering,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 37, no. 5, pp. 2680–2694, 2025.
- [8] J. Liu and S. Chen, “Timesurl: Self-supervised contrastive learning for universal time series representation learning,” in *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, vol. 38, no. 12, 2024, pp. 13 918–13 926.
- [9] S. Gao, T. Koker, O. Queen, T. Hartvigsen, T. Tsiligkaridis, and M. Zitnik, “Units: A unified multi-task time series model,” in *Proceedings of the Annual Conference on Neural Information Processing Systems (NeurIPS)*, vol. 37, 2024, pp. 140 589–140 631.